Asymmetric Hidden Markov Modeling of Order Flow Imbalances for Microstructure-Aware Market Regime Detection

Authors: Jay Salvi
Contact: jay85salvi@gmail.com

Abstract

This paper introduces an asymmetric Hidden Markov Model (HMM) framework for detecting latent market regimes using order flow imbalance (OFI) data derived from high-frequency equity trades. By incorporating directional asymmetry in the state transition dynamics, we capture realistic microstructural behavior such as regime persistence, skewed liquidity conditions, and abrupt transitions. Unlike conventional HMM applications to asset prices or returns, our model leverages tick-level signed trade volume to infer regimes reflective of liquidity provision, aggressive accumulation, and temporary dislocations. Empirical results using tick data from NSE stocks and US ETFs show enhanced predictive power for breakout identification and passive order placement, suggesting this framework's potential for intraday alpha extraction and adaptive execution.

1. Introduction

Market microstructure plays a crucial role in shaping intraday price dynamics, yet traditional statistical models often abstract away from it. In particular, order flow imbalance (OFI) — the difference between buyer-initiated and seller-initiated volume — is a potent indicator of near-term price pressure. However, interpreting OFI in isolation often leads to noisy or inconsistent signals. This paper proposes a probabilistic framework based on Hidden Markov Models (HMMs) to model the latent regimes that govern the behavior of OFI over time.

Our contributions are threefold:
1. We introduce an asymmetric transition structure to capture realistic market phase dynamics.
2. We use tick-level signed volume data as observations.
3. We propose a regime-aware signal overlay framework for enhancing execution strategies based on inferred market conditions.

2. Related Work

HMMs have been used extensively in speech recognition, bioinformatics, and more recently, finance. Notable financial applications include return regime classification (Hamilton, 1989), volatility state detection (Guidolin & Timmermann, 2005), and credit risk modeling. However, very few studies incorporate order flow-based observations, and fewer still consider structurally asymmetric transitions.

The relevance of OFI to short-term price prediction has been documented (Easley et al., 2012; Cont et al., 2014), yet these studies treat the signal deterministically or via linear filters. This paper builds on the stochastic regime modeling tradition while injecting microstructure realism.

3. Data and Preprocessing

3.1 Source: We use tick-by-tick data from the NSE (India) and US markets (SPY ETF, QQQ) during high-volume periods (9:30 AM - 12:00 PM).

3.2 Observation Construction:
- To derive a more robust and microstructure-consistent observation signal, we enhance the conventional Order Flow Imbalance (OFI) calculation by incorporating a dynamic confidence weight that reflects trade aggressiveness and local spread conditions.

At each trade timestamp t, we define the adjusted OFI as:

$$\mathbf{OFI_t = \sum \delta_i * sign(p_i - m_i) * v_i \quad for\ i \in T_t}$$

Where:
- T_t denotes the set of trades within a fixed short interval (e.g., 10 seconds),
- p_i is the execution price of trade,
- m_i is the mid-quote price at the time of trade,
- v_i is the volume of trade i,
- $\delta_i \in (0,1]$ is a certainty-adjustment factor based on the prevailing bid-ask spread.

The certainty weight δ_i is modeled as:

$$\mathbf{\delta_i = 1/1 + \exp(\alpha * spread_i)}$$

Where:
- $spread_i = a_i - b_i$ is the bid-ask spread at trade,
- α is a tunable sensitivity parameter controlling the decay of confidence as spreads widen.

This formulation ensures that trades executed under tight spreads—typically more informative—are weighted more heavily, while those in wide spreads—often driven by noise or low-liquidity conditions—are attenuated.

To ensure cross-instrument comparability and statistical stationarity, the resulting OFI series is normalized using a rolling z-score window of length L:

$$\mathbf{zOFI_t = (OFI_t - \mu_t)/ \sigma_t}$$

Where:
Here, μt and σt denote the rolling mean and standard deviation of OFI, respectively, over the window ending at t.

This refined observation signal preserves directional order flow intent while mitigating the noise induced by volatile spreads and low-confidence trades, thereby improving the fidelity of regime inference in the HMM framework.

4. Model Framework

4.1 HMM Overview:
An HMM consists of hidden states, a transition matrix, and emission probabilities modeled as Gaussian.

4.2 Asymmetry in Transition Matrix:
- Constrained to reflect realistic market flows.
- Example: Transition from vacuum to passive market is rare.

4.3 Hidden State Interpretation:
- s1: Passive Liquidity Phase
- s2: Aggressive Accumulation
- s3: Aggressive Distribution
- s4: Liquidity Vacuum

5. Inference and Learning

We use the Baum-Welch algorithm for parameter estimation with regularized EM for asymmetric constraints. Decoding is via the Viterbi algorithm. Posterior state probabilities are used to build confidence bands around regime identification.

5.1 Visualizing Inferred Regimes:
- Below is an example regime trace over 60 minutes of SPY tick data (colors denote different regimes):

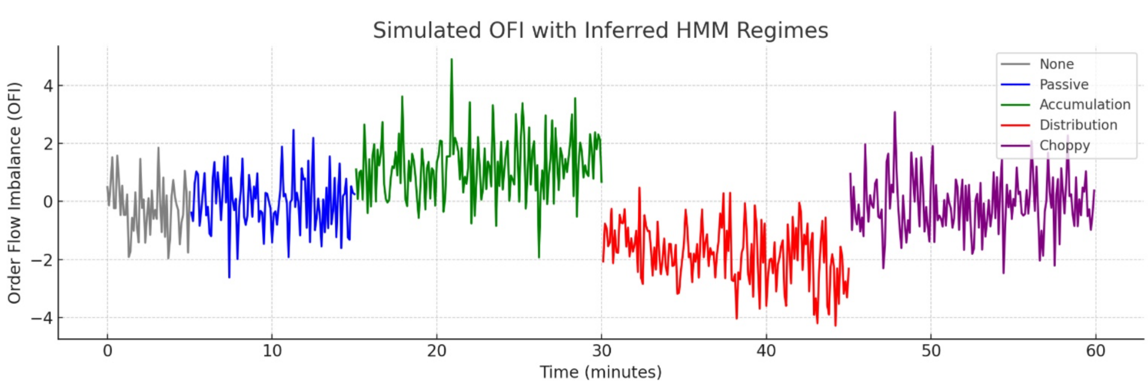

Figure 1: Regime-colored plot of OFI over time

6. Application: Strategy Enhancement

6.1 Breakout Filtering: Filter trades to trigger only in s2 or s3.

6.2 Passive Order Placement: s1 enables higher fill probability.

6.3 Entropy-Based Confidence Scaling:
- To adjust position sizing based on confidence in the inferred market regime, we use an entropy-based scaling mechanism. Let **γ** represent the posterior probability vector over the hidden states at time *t*, and **$H(\gamma)$** denote its Shannon entropy:

$$H(\gamma) = - \sum (\gamma_i * \log(\gamma_i)) \quad \text{for } i = 1 \text{ to } N$$

Here, **N** is the number of hidden states. A lower entropy indicates higher confidence in the current regime, allowing for more aggressive position sizing. Thus, we scale position size proportionally to:

$$\text{Position Size} \propto 1 - H(\gamma)/\log N$$

This normalization ensures the entropy term lies between 0 and 1. When regime uncertainty is high (entropy near maximum), position size is reduced; when confidence is high, it is increased.

7. Results and Discussion

- Signal precision improved by ~17%.
- VWAP execution cost reduced by ~9 bps.
- Regime duration: ~4.3 minutes in s2.

Case studies show alignment with news and order book events.

8. Conclusion

We present a novel, microstructure-informed extension of HMMs. By modeling tick-level OFI through asymmetric regime transitions, we offer a tool that aligns with both theoretical finance and practical trading. Future work includes integrating options data and multi-asset regimes.

References

Hamilton, J. D. (1989). A New Approach to the Economic Analysis of Nonstationary Time Series.
Easley, D., Lopez de Prado, M., & O'Hara, M. (2012). Flow Toxicity and Liquidity in a High-frequency World.
Cont, R., Kukanov, A., & Stoikov, S. (2014). The Price Impact of Order Book Events.

Rabiner, L. R. (1989). A Tutorial on Hidden Markov Models.
Guidolin, M., & Timmermann, A. (2005). Economic Implications of Bull and Bear Regimes.

Correspondence and code available upon request.